Linear Regression

In [1]:
```
import pandas as pd
import numpy as np
```

In [2]:
```
import matplotlib.pyplot as plt
import seaborn as sns
```

In [3]:
```
%matplotlib inline
```

In [4]:
```
df = pd.read_csv('USA_Housing.csv')
```

In [6]:
```
df.head()
```

Out[6]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanielton, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386

In [7]:
```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
 #   Column                        Non-Null Count  Dtype  
---  ------                        --------------  -----  
 0   Avg. Area Income              5000 non-null   float64
 1   Avg. Area House Age           5000 non-null   float64
 2   Avg. Area Number of Rooms     5000 non-null   float64
 3   Avg. Area Number of Bedrooms  5000 non-null   float64
 4   Area Population               5000 non-null   float64
 5   Price                         5000 non-null   float64
 6   Address                       5000 non-null   object 
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

In [8]:
```
df.describe()
```

Out[8]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

In [9]:
```
df.columns
```

Out[9]:
```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

In [10]:
```
sns.pairplot(df)
```

Out[10]:
```
<seaborn.axisgrid.PairGrid at 0x17870c739a0>
```

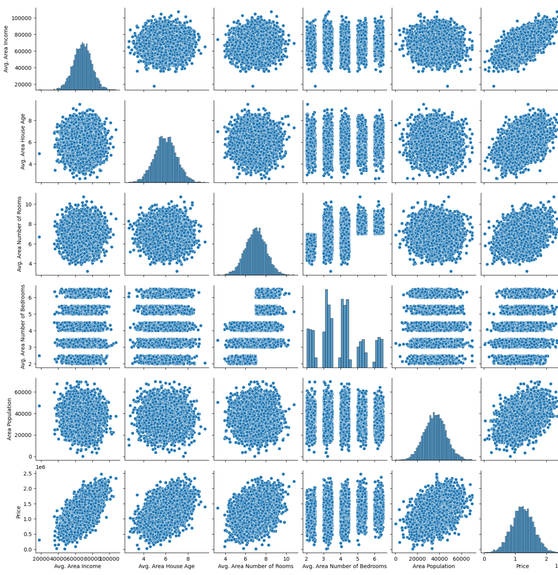

In [12]:
```
sns.distplot(df['Price'])
```

```
C:\Users\tsaks\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[12]:
```
<AxesSubplot:xlabel='Price', ylabel='Density'>
```

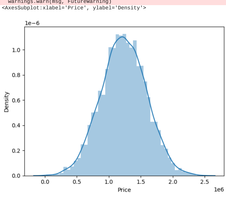

In [14]:
```
sns.heatmap(df.corr(), annot = True)
```

Out[14]:
```
<AxesSubplot:>
```

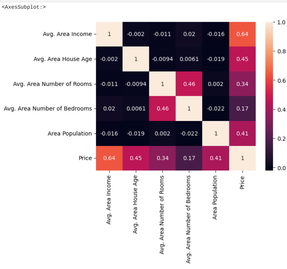

In [15]:
```
df.columns
```

Out[15]:
```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

In [16]:
```
X = df[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population']]
```

In [17]:
```
y = df['Price']
```

In [18]:
```
from sklearn.model_selection import train_test_split
```

In [19]:
```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=101)
```

In [20]:
```
from sklearn.linear_model import LinearRegression
```

In [21]:
```
lm = LinearRegression()
```

In [22]:
```
lm.fit(X_train,y_train)
```

Out[22]:
```
LinearRegression()
```

In [23]:
```
print(lm.intercept_)
```

```
-2640159.79685257267
```

In [24]:
```
lm.coef_
```

Out[24]:
```
array([2.15282755e+01, 1.64883282e+05, 1.22360678e+05, 2.23380186e+03,
       1.51594209e+01])
```

In [25]:
```
X_train.columns
```

Out[25]:
```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population'],
      dtype='object')
```

In [27]:
```
cdf = pd.DataFrame(lm.coef_,X.columns, columns=['Coeff'])
```

In [28]:
```
cdf
```

Out[28]:

	Coeff
Avg. Area Income	21.528276
Avg. Area House Age	164883.282027
Avg. Area Number of Rooms	122360.678027
Avg. Area Number of Bedrooms	2233.801864
Area Population	15.150420

Predictions

In [29]:
```
predictions = lm.predict(X_test)
```

In [30]:
```
predictions
```

Out[30]:
```
array([1260960.70567827,  827588.75566334, 1742421.2425434 , ...,
        372191.40628023, 1365217.15140897, 1914519.5417887 ])
```

In [31]:
```
plt.scatter(predictions,y_test)
```

Out[31]:
```
<matplotlib.collections.PathCollection at 0x178765ef370>
```

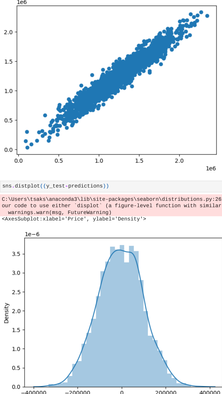

In [32]:
```
sns.distplot((y_test-predictions))
```

```
C:\Users\tsaks\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[32]:
```
<AxesSubplot:xlabel='Price', ylabel='Density'>
```


In [33]:
```
from sklearn import metrics
```

In [34]:
```
metrics.mean_absolute_error(y_test,predictions)
```

Out[34]:
```
82288.22251914957
```

In [35]:
```
metrics.mean_squared_error(y_test,predictions)
```

Out[35]:
```
10460958907.209059
```

In [36]:
```
np.sqrt(metrics.mean_squared_error(y_test,predictions))
```

Out[36]:
```
102278.82922290936
```

In [ ]:
```

```